

Harsh Bhanushali

Agentic AI Developer — Multi-Agent Systems and RAG Pipelines
harshbhanu0709@gmail.com | 7041132742 | Vapi, Gujarat | [LinkedIn](#) | [GitHub](#) | [Portfolio](#)

Summary

Agentic AI Developer building multi-agent systems and RAG pipelines in Python — architecture-first, with full control over retrieval, routing, and agent state.

Skills

AI / ML	Multi-Agent Systems, RAG, LLM Integration, Tool-Calling, Prompt Engineering, Function Calling, Embeddings, Vector Databases, Groq, Gemini, OpenAI-compatible APIs
Languages & Libraries	Python, SQL, NumPy, Pandas, Pydantic, LangChain, LangGraph, SentenceTransformer
Frameworks & Tools	FastAPI, Streamlit, ChromaDB, MCP (Model Context Protocol)

Experience

Symbiosis Systems and Services Pvt. Ltd.

Jan 2024 – Sep 2024

Full Stack Developer Intern

- Delivered 20+ features across 15+ backend modules in 3 enterprise web applications using ASP.NET and SQL Server; reduced manual workflow steps by 30% for daily active users.
- Spearheaded the architectural overhaul of the core enterprise application, rearchitecting the codebase to engineer automated, cross-module data workflows that eliminated system friction.

Projects

AI Agent Engine github.com/harshbhanushali26/ai-agent-engine

Feb 2026 – Present

- 4-layer pipeline (cache → pattern matcher → RAG → LLM) — ~80% of queries never reach the LLM, keeping inference cost near zero for repeat usage patterns.
- Five sequential agents (Planner → Validator → Executor → Replanner → Responder) with auto-retry and replanning on failure. Every decision logged at session level for full traceability.

Stack: *Python, LangChain, ChromaDB, Gemini, FastAPI, Pydantic*

hArI — Document Intelligence App harichat.streamlit.app

Mar 2026 – Present

- Rebuilt retrieval on Supabase — hybrid search combining vector similarity and keyword ranking; replaced pandas exec() with in-memory DuckDB enabling secure multi-CSV JOIN queries via LLM-generated SQL.
- Ships with Supabase Auth, per-user chat persistence, and a citation UI that shows the exact retrieved chunks for each answer.

Stack: *Python, Streamlit, ChromaDB, Groq, LangChain, SentenceTransformer*

Multi-Agent Research Pipeline github.com/harshbhanushali26/research-pipeline

Apr 2026 – Present

- 6 agents each own one stage — search, scrape, summarise, critique, synthesise. Supervisor routes and triggers re-search when gaps are found in earlier outputs.
- Automatically detects research gaps mid-run and re-searches before synthesising — final report covers 30–50 sources across 5–9 targeted questions per run.

Stack: *Python, Groq, Gemini, DuckDuckGo API*

NextSteps — Career Path AI github.com/harshbhanushali26/NextSteps

Mar 2026 – Present

- Full resume-vs-JD analysis in 5–7 min end-to-end via FastAPI — returns gap analysis, skill map, and learning plan.
- Schema-agnostic resume parser and dynamic JD extractor — no hardcoded field assumptions, works on any format or live job URL.

Stack: *Python, FastAPI, Groq, Gemini, Pydantic, REST API*

DevMind — MCP Server github.com/harshbhanushali26/mcp-implementations

May 2026 – Present

- Local MCP server with 6 tools — file read/write, Python execution, JSON utilities — accessible to any MCP-compatible client over stdio.
- HITL confirmation on all write and run actions. Fixed a subprocess stdin isolation bug that silently hangs most naive MCP implementations.

Stack: *Python, MCP (Model Context Protocol), subprocess, HITL*

Education

Government Engineering College, Modasa — B.E. Information Technology

Sep 2020 – May 2024

CGPA: 8.23 | GATE 2025 CSE — Qualified